

Execution

Purpose: give a detailed account of the build of the project.

Introduction

The following chapter gives an outline of the process that was undertaken to build out the final application. A description of early prototype work is given to give context to the construction of the final working prototype version. This is followed by a description of the technical architecture of the system as well as an outline of the sometimes tricky setup process of getting the live reload system working (as described in the previous chapter). Some detailed discussion of the the core functionality of the system is then given. The core of the system primarily consists of timeline events, which visually manifest themselves as strokes on the canvas and aurally as fm synthesized frequency modulated sounds. As will be discussed an important aspect of any well constructed software system is a good degree of separation of concerns, a characteristic espoused to and evident in the resulting code. To this end, the code that brings the central functionality to life can conceptually be divided into a data or entity layer, a business logic or use case layer and an output layer which in this case consists of the visual output into a html canvas element and the audio output through the web audio api. In this sense the architecture conforms to the principles of the Clean Architecture as presented by Bob Martin (??).

Early prototype work

Melodypainter

Melodypainter is an early prototype built out in Max MSP that allows users to draw freehand lines, which are converted into break point function data and used to generate a melodic profile using Bach for Max MSP. Bach is a suite of composition tools that allow for a number of computer aided composition techniques (CAC) and provides similar functionality to IRCAM's Open Music system. These melodic profiles are then filtered to only include notes from a pentatonic scale, to give reasonably pleasing aural results. Some notable flaws in the system include the following. It is limited to strictly western tonal music styles. It has no allowance for rhythm and plays only eight notes giving results a noticeably bland and predictable quality. The freeform nature of sketched input however was quite a pleasing means of inputting the control information.

Sonicshaper [0/1]

- [] Add a picture

A separate application was created in Processing which allowed users to draw shapes, using either mouse or ideally, pen input and have a sound that is associated with each shape played back. As the sound of each shape plays back, it is lit up using animation, creating a strong connection between the shape and it's resulting sound. The application uses the “gesture variation follower” system [caramiaux_adaptive_2015], which while promising in principle, didn't have a high rate of accuracy in recognizing the shapes.

TODO Web version of William Coleman's SonicPainter [0/1]

- [] Reference processing.js
- OSC
- Tone.js
- WC SonicPainter
- Microsoft calculator - Harel Statecharts

A potential starting point that was considered was using the code from William's SonicPainter and porting it to the web platform. This process proved to be quite straightforward. The processing code could more or less be embedded in a webpage as is using “processing.js”, a web version of the Processing library that enables users to run processing sketches in the Web Browser. Some notable changes that had to be made were removing the OSC functionality as this is not technically possible to use in a browser. In addition, some other pieces of code had to be commented out and tweaked. As it's not possible to run Max MSP patches in the browser, the audio system was re-implemented using Tone.js. As SonicPainter uses simple FM synthesis, a very close approximation to the original version could be created. In the end, it was decided not to build on this codebase however as there were some issues with functionality and useability that would be difficult to resolve in an inherited codebase. A fundamental issue was that transitions between certain states would cause crashes or unpredictable behaviour. An example of this is when a user attempts to use the vibrato tool while in the process of creating a note. Instead of either finishing the note and starting the vibrato tool or disallowing the behaviour, the program would crash. This is a common problem in software development and is evidence even in commercial products. To alleviate such issues in the new codebase a more disciplined approach would be taken to managing transitions between states. The process of porting the code did however give a more in depth incite into Coleman's implementation including the pitfalls mentioned above. In addition, the basic visual look and conceptual functionality would form the basis of the workings of SonicSketch.

Actual implementation

Setting up the architecture

1. Clojurescript and javascript npm modules

- [] Ref clojurescript age
- [] Ref new innovations with clojurescript
- NPM
- Node.js
- Project.clj
- Clojurescript

Despite the fact that clojurescript has existed for six years(???), some areas of the development process are still difficult, particularly when building out a more complex real world application. It should be noted that a good deal of work is being carried out to make this a smoother experience and thusly these pains are likely to become less of an issue in the near future (???ref). It should also be noted that building applicatons using plain javascript is not a trivial process and in all likelihood will include a build process using a system like webpack or browserify. A primary issue that had to be resolved to allow the application to be built out was the incorporation of javascript npm modules. NPM is the module system used by node.js originally for more server oriented technologies but increasingly for rich clientside applications. For a purely javascript application, it would be a matter of simply adding the desired libraries as dependencies. However, with the use of clojurescript some extra steps needed to be carried out. In addition to adding the dependencies, a javascript file was created that imported these into a “deps” object. This deps object could then be referred to in clojurescript using the standard interop syntax `js/deps.myDependency`. At the time of development an alpha feature that allowed npm dependencies to be declared as part of the project.clj file was experimented with but was not used in due to difficulties getting it to work. While the project setup was not as elegant or succinct as might be wished, it did provide a stable base to build on and a means to harness the rich resource that is the NPM ecosystem and use such tools as Paper.js and React.js.

2. Paper.js and react.js (paper.js bindings)

- Scenegraph
- Binding

As has been outlined in the previous chapter, a declarative coding style would be employed to enable a live coding workflow and to avoid a building a codebase that is increasingly difficult to understand. In other words the code should as much as possible describe the “what” of the functionality

rather than the “how”. These qualities emerge quite naturally when using the *React.js* architecture. Paper.js however runs in the context of a canvas element and thusly it is not possible to directly use *React.js* with it. This shortcoming has been addressed in projects such as *three.js react bindings* and *pixi.js react bindings* which allow the use of react’s declarative programming style for 3d and 2d scenegraph oriented systems that run in the html canvas element. These solutions both work by creating dummy empty dom elements and hook into the *React.js* lifecycle events to the real work of updating the scenegraph. In many ways the scene graph structure of projects like these and indeed Paper.js exhibit a high resemblance to DOM structures and APIs making React a good fit for them. A similar approach to the above mentioned libraries approach was taken to integrate paper.js for use in SonicSketch and worked reasonably well but required quite a bit of setup. During the development of the project, a more suitable solution emerged from the open source community at an opportune time. This used the next version of *React.js* which has better support for render targets that are not the DOM and has the distinct advantage of not requiring the creation of redundant DOM nodes. The library was far from comprehensive and thusly a custom version of the library was used that included some custom functionality required for SonicSketch.

3. Tone.js and react.js

In some ways audio output can be thought of in a similar way to the visual output of the app and thusly can be treated in similar way by *React.js*. It can use the declarative data oriented system of react to configure the particular settings and connections in the audio graph and hook in to its lifecycle events to instantiate the various audio generating and processing web audio nodes. This addresses a notable (by design) omission in Tone.js which does not allow the code to query the state of the audio graph once it has been setup. It is down to the userland code to keep track of this and manage it accordingly. The value proposal offered by introducing react.js into this part of the system is that it maintains the simple relationship between state and generated output. Conceptually the flow of change is:

- (a) The state updates
- (b) The react wrapper objects update their properties accordingly
- (c) The lifecycle events are triggered which takes care of altering, adding and removing web audio nodes (thus altering the audio being output)

The design of this part of the application is influenced by *react music*, a system that uses *React.js* with *tuna.js*, a web audio library similar to *tone.js* (???ref).

4. Reagent and react.js paper.js bindings

The final piece of the jigsaw in the underlying technology stack is the integration of react with clojurescript via the *Reagent* library. The core syntax of this system is simple clojurescript vectors similar to the following:

```
[:div
  "Hello " [:span {:style {:font-weight bold}}
world]]
```

This would result in the following html output:

```
<div>Hello <span style="font-weight: bold">world</span></div>
```

As can be seen the vectors begin with a keyword that corresponds to the tagname of the html. Additionally, instead of using html tag keywords, function calls can be made to generate html to allow for code reuse and logic. It was unclear how the paper.js bindings would work within this system due to the fact that it required a different version of react and uses non standard tag names for elements that can be drawn on screen such as “circle” and “rectangle”. This however turned out to be much more straightforward than expected and the provided paper.js primitives could be simply using the relevant paper.js keywords. Complex scenegraphs could be constructed by using the following succinct clojurescript syntax to describe the playback indicator:

```
[:Group {:position [position 0]
          :pivot [0 0]
          :opacity 0.75}
 [:Rectangle {:pivot [0 0]
              :size [1 height]
              :fill-color "#ffffff"}]
 [:Path {:segments [[[-5 0] [5 0] [0 7] [-5 0]]
                    :fill-color "#ffffff"}]]
```

As can probably be inferred from the code the `position` and `height` are properties that are passed into the hiccup and trigger updates to the visual display when they change: in the case of position, when the playback position changes and in the case of the height, when the user resizes the browser window. The path element describes the triangle that is placed at the top of the screen.

The current state of the art in live code reloading in the browser is still not as comprehensive or as easy to setup as might be wished. Once it has been configured it is difficult to return to the compile and run workflow, and is in most cases a worthwhile investment of time. With these underlying elements in place the process of creating the core functionality application could begin and will now be described.

Core functionality - timeline events (or notes)

1. Introduction
 - Describe the core functionality
 - Describe core entities

2. Add timeline event
 - Business logic
 - UI
 - Audio
3. Add vibrato
 - Business logic
 - UI
 - Audio
4. Remove note
5. Move note
6. Change sound (preset system)
7. Probability

At the core of the application is the creation of timeline events which unfold in a looped fashion. These events are created based on the input of the user with a mouse or mouse like input device. On the production of a valid input gesture, the screen is updated immediately with a visual display of this content. The details of this gesture is stored in memory and the event that will eventually create the sound is registered with tone.js. Much of the events that occur in the system are captured in a main “View” component which houses the central html canvas element. To aid in organising the large amount of functionality associated with the component, higher order components are used to separate this out into logical groupings. A higher order component is a component that wraps a normal component to add functionality to it and accepts the same properties as the component it wraps (???ref-react). In this case the most logical grouping is by tool and so there are higher order components setup for each of the tools: draw, vibrato, delete, move, resize and probability.

1. Add timeline event
 - Clojure atom

This is the default tool that is activated when the user opens the application and enables the user to add timeline events by drawing them onto the screen. It is added to the system when the “draw” tool is activated and a mouse drag operation is carried out from left to right within the bounds of the canvas element. The event is captured in the main canvas view and is initiated when the user left clicks the mouse triggering the following function:

```
(defn pointer-down [{:keys [temp-obj active-preset]} evt]
  (let [pointer-point (.. evt -point)
        group        (js/paper.Group. (clj->js {:position
                                                    [(.. pointer-point -x) (.. pointer-point -y)]
                                                    :applyMatrix false
                                                    :pivot [0 0]}))
        circle        (js/paper.Shape.Circle. (clj->js {:fillColor "#ffffff"
                                                         :radius      5}))]
```

```

    path          (js/paper.Path. (clj->js {:strokeColor "#ffffff"
                                              :strokeWidth 2
                                              :fullySelected true
                                              :segments [[0 0]]}))

  (.. group (addChildren #js [circle path]))
  (reset! temp-obj {:path path
                    :circle circle
                    :group group
                    :loc pointer-point}))

```

This function receives a hashmap with a reference to a *ClojureScript* atom which can be `reset!` to contain a temporary visualisation of the newly created note. This function uses a heavy amount of javascript interop to directly instantiate paper.js objects and add them to a shared group.

As the user continues to move the cursor further points are added to the path created in the `pointer-down` function. Some constraints however are placed on the creation of the path and only points that are past the last previous from left to right are added. If the user backtracks it leads to a deletion of points, providing an intuitive undo like behaviour and implementing a reciprocal HCI interaction pattern recommended by NUI principles.

```

(defn pointer-move [{:keys [temp-obj active-preset]} evt]
  (when-let [{:keys [path group] :as temp-obj} @temp-obj]
    (let [pointer-point (.. evt -point)
          rel-pos       (.. group (globalToLocal pointer-point))]
      ;; Only add positive points relative to first
      ;; Remove points greater than pointer-points
      (when-let [last-seg (.. path getLastSegment)]
        (let [first-seg (.. path getFirstSegment)
              first-point (-> first-seg .-point)
              last-point  (-> last-seg .-point)
              pointer-x    (.-x rel-pos)
              amp-env      (-> active-preset :envelope)
              stage-width  (.. evt -tool -view -viewSize -width)
              max-width    (if (= (-> amp-env :sustain) 0)
                            (let [time (+ (-> amp-env :attack)
                                           (-> amp-env :decay)
                                           (-> amp-env :release))]
                              (-> time
                                ;; Seconds to beats
                                (* (/ js/Tone.Transport.bpm.value 60))
                                (time->euclidian stage-width)))
                            nil)]
          (when (or
                 (nil? max-width)

```

```

        (< pointer-x (+ (-x first-point) max-width)))
    (-> path (.add rel-pos))
    (let [greater-segs (filter
                        #(> (-> % .-point .-x) pointer-x)
                        (-segments path))]
        ;; Remove greater points
        (doseq [seg greater-segs]
            (.removeSegment path (-index seg)))))))))

```

Completion of a note occurs when the user releases the button and triggers the `pointer-up` function:

```

(defn pointer-up [{:keys [temp-obj active-preset stage-size]} evt]
  (let [{:keys [path circle group loc] :as temp-obj} @temp-obj]
    (.simplify path 10)
    ;; Send the actual note
    (dispatch [:note-add (-> (path->note path loc stage-size)
                              (assoc ,,, :preset active-preset)
                              ;; Use the color from the active preset
                              (assoc ,,, :color (:color active-preset)))] )

    ;; Remove temp stuff
    (.remove path)
    (.remove group)
    (.remove circle))
    ;; Unset temp obj
    (reset! temp-obj nil))

```

This function simplifies the path by calling the `paper.js simplify` method on the path object and dramatically reduces the amount of data captured while preserving the basic characteristic of the user's stroke (???`ref-simplify-fn`). Most importantly it calls the *re-frame* `dispatch` function to add the note to the app database. A `path->note` function is used to convert the stroke from the domain of euclidean space on the visual space of the canvas to the domain of time-pitch space for use with the audio synthesis system. The `path->note` function can be seen below:

```

(defn path->note [path first-point stage-size]
  "Main entry point to this namespace"
  (let [path-width (.. path -bounds -width)
        width      (:width stage-size)
        height     (:height stage-size)]
    {:freq      (domain/euclidean->freq (.. first-point -y) height)
     :onset     (domain/euclidean->time (.. first-point -x) width)
     :duration  (domain/euclidean->time path-width width)
     :velocity  0.5
     :enabled   true
     :probability 1.0
     :color     @(col/as-css (get colors (rand-int 100)))})

```



```

:height      (... path -bounds -height)
:width       (... path -bounds -width)
:envelopes   {:frequency {:raw      (paper-path->vec path [width height])
                           :sampled (paper-path->sample path stage-size)}
               :vibrato   (reduce (fn [a b] (assoc a b [b 0])) (sorted-map) (range

```

The domain of time-pitch is used to store the notes in memory and makes it possible to maintain a relative relationship between the screen size and the drawn notes.

The dispatched note event is handled by a *re-frame* `reg-event-db` handler which describes the alteration that is required to be made to the database. It also uses a series of interceptors, to perform validation of the database and to remove some of the boilerplate from the event handler functions. Interceptors are similar conceptually to middleware and is the place where all of the side effects arising from an event are actioned. Moving application side effects to this placed ensures that they are isolated and means that the majority of the program can be kept as pure logic improving to make it easier to test, debug and reason about. As can be seen the handler function is very simple:

```

(reg-event-db
 :note-add
 note-interceptors
 (fn [notes [note-info]]
  (let [id      (allocate-next-id notes)
        note    (assoc note-info :id id)]
    (if (>= (:duration note) 0.001)
      (assoc notes id note)
      notes))))

```

This does a simple check to make sure that note has a minimum duration and if so alters the notes vector to include the new vector which will instruct *re-frame* to update it's internal atom with this new state.

The structure of the note hashmap is defined using *clojure.spec*, a core library to perform data validation and a tool that may be used similarly to types in strongly typed languages. The note specs are defined as follows:

```

(s/def ::id int?)
(s/def ::freq float?)
(s/def ::onset float?)
(s/def ::duration float?)
(s/def ::velocity float?)
(s/def ::color string?)
(s/def ::note (s/keys :req-un [::id ::freq ::onset ::duration]
                      :opt-un [::velocity]))

```

Although note specified here, notes also have an `envelopes` key that stores

frequency and vibrato envelopes. For example:

```
{:frequency {:raw [{:point [-0.01 99.7] :handle-in [0 100] :handle-out [0 99.8]} ...  
                  {:point [2.8 98.02] :handle-in [-0.10 100.4], :handle-out [0 100]}]  
              :sampled [-0.36991368680641185 ... -2.172026174596564]}  
:vibrato {0 [0 0], ... 10 [10 0]}}
```

The update in state spins *re-frame*'s subscription system into action and any views that are subscribed to the application state are now re-rendered as needs be. The `graphics-notes` view for instance is subscribed to `:notes`, `:graphics-stage-width`, `:graphics-stage-height`: and `:mode`:

```
(defn graphics-notes []  
  (let [notes      (subscribe [:notes])  
        width      (subscribe [:graphics-stage-width])  
        height     (subscribe [:graphics-stage-height])  
        ....]  
    ;; playback-time @(subscribe [:playback-time])  
    (into [:Group]  
      (map (fn [note]  
             ^{:key (:id note)} [graphics-note* ...]) @notes))))
```

When a note is added this render function will be ran which will call the `[graphics-note*]` component for each of the notes, and update the visual display of the screen to show the new note. A similar process happens in the audio system except in this case, new web audio nodes are created, timeline events are queued up at the correct time and audio envelopes are setup that trace the curve of the drawn lines.

```
(defn fm-synth [{:keys [out] :as props} & children]  
  "FM synth"  
  ;; The new synth is instanciated using js interop  
  (let [synth (js/Tone.FMSynth. (clj->js (dissoc props :out)))]  
    (reagent/create-class  
      {:component-did-mount  
       (fn []  
         ;; The synth is connected to it's output (which is passed  
         ;; in as a property to the component)  
         (.. synth (connect out)))  
       :reagent-render  
       (fn [props & children]  
         ;; The render function renders a dummy span dom element and  
         ;; renders it's children and passing it's synth as the out  
         ;; for these components.  
         (into [:span]  
           (map (fn [child]  
                  (assoc-in child [1 :out] synth))  
                children))))
```

```

:component-will-unmount
(fn []
  ;; Here we dispose of the synth
  ;; This will happen when the parent note is removed or when
  ;; a live code reload happens
  (... synth dispose))))))

```

The above shows the `fm-synth` which sits at the heart of the audio generating part of the system. Potential parent components of this would be audio effects, or the master bus. It's child components are comprised of note events and envelopes that drive frequency changes over the course of the note playback. The composition of events, envelopes, synths, effects and channels is shown in truncated form:

```

;; Parent component is the project and
;; has settings such as tempo
[project {:project :settings}
 [master-bus {}
  ;; Master volume is set here
  [volume {:volume :settings}
   ;; Adds a simple reverb effect
   [reverb-effect {}
    [
      ;; First note
      ;; Each note has a vibrato effect
      [vibrato-effect {}
       ;; Envelope to control vibrato depth
       [timeline-evt evt
        [envelope {:param "depth"
                   :env vib-env}]]
       ;; The fm synth that generates the signal
       [fm-synth (get-in evt [:preset])
        ;; Timeline event that takes care of queuing
        ;; it's child components
        [timeline-evt evt
         ;; In this case a note
         [note {:note :settings}]
         ;; And a frequency envelope
         [envelope {:param "frequency"
                    :state state
                    :env freq-env}]]]]
      ;; Second note
      [vibrato-effect {}
       ;; ...
      ]
      ;; ...
    ]]]]]

```

2. Editing notes

- Vibrato
- Remove note
- Move note
- Resize note
- Change sound (preset system)
- Probability tool

The

Secondary functionality

1. Introduction
2. Transport controls
3. Animation (current play position & notes)
4. Undo and redo
5. Fullscreen
6. Outer UI
7. Save and load file

Performance issues

Conclusion

- Summarise the resulting artifact